



Natural Language Processing

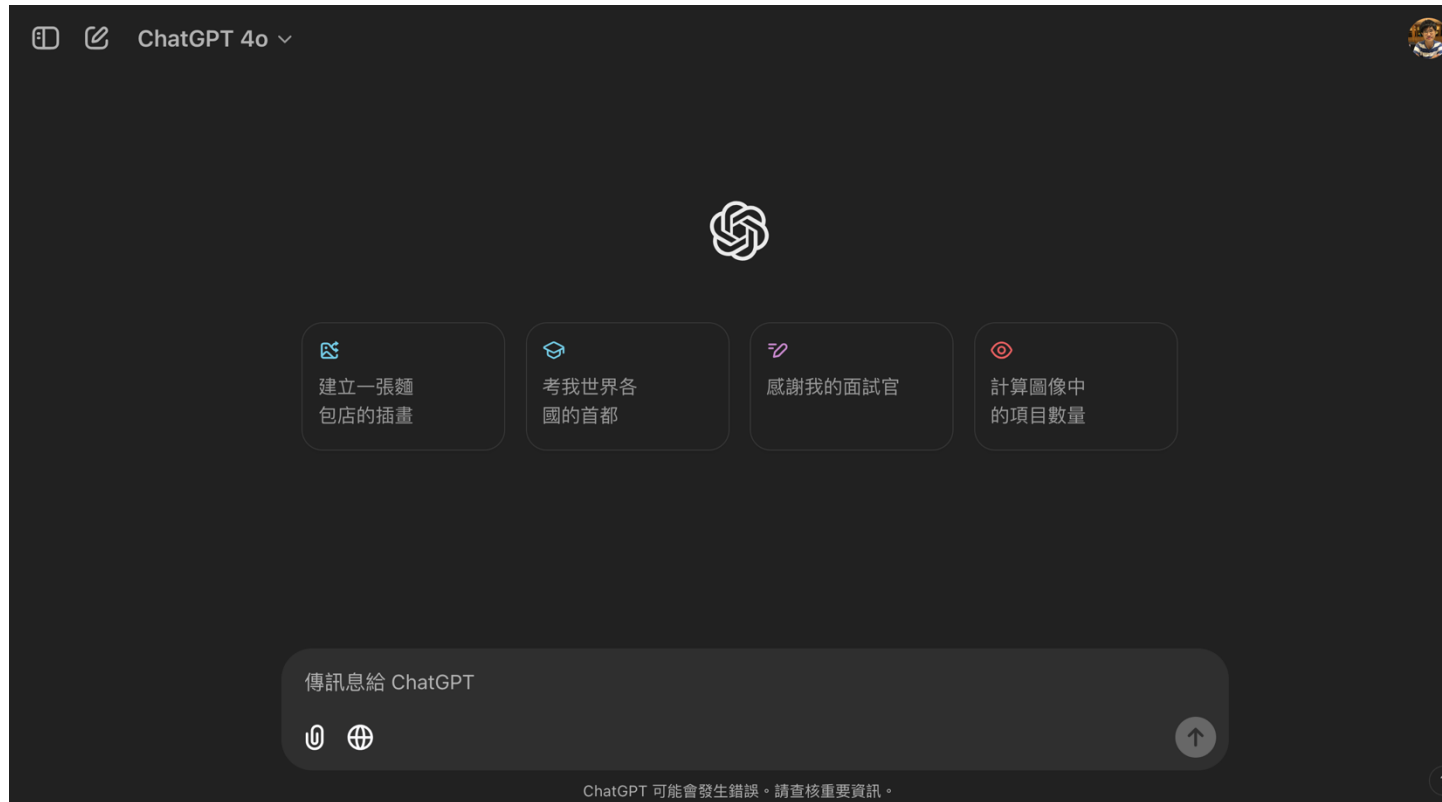
LLM API Tutorial 2024/11/21



What you will learn in this tutorial

- Gemini API -> **Google** Gemini
- OpenAI API -> **OpenAI** ChatGPT
- Claude API - > **Anthropic** Claude
- Code (notebook) Link
 - https://github.com/IKMLab/NTHU_Natural_Language_Processing/tree/main/Reference/LLM_API_lab

Why do we need to use API?



- It's slow if you want to use ChatGPT for NLP tasks by manually copying and pasting from the webpage.
- If you use ChatGPT from the webpage to test data, "Too many requests in 1 hour. Try again later." may occur.

How to choose a good LLM?

Chatbot Arena LLM Leaderboard: <https://lmarena.ai/?leaderboard>

Rank* (UB)	Rank (StyleCtrl)	Model	Arena Score	95% CI	Votes	Organization	License
1	4	Gemini-Exp-1114	1344	+7/-7	6446	Google	Proprietary
1	1	ChatGPT-4o-latest (2024-09-03)	1340	+3/-3	42225	OpenAI	Proprietary
3	1	o1-preview	1333	+4/-4	26268	OpenAI	Proprietary
4	5	o1-mini	1308	+4/-3	28953	OpenAI	Proprietary
4	4	Gemini-1.5-Pro-002	1301	+4/-4	23856	Google	Proprietary
6	9	Grok-2-08-13	1290	+3/-3	47908	xAI	Proprietary
6	11	Yi-Lightning	1287	+4/-4	27114	01 AI	Proprietary
7	4	GPT-4o-2024-05-13	1285	+2/-2	108575	OpenAI	Proprietary
7	3	Claude 3.5 Sonnet (20241022)	1283	+4/-4	26047	Anthropic	Proprietary
10	16	GLM-4-Plus	1275	+3/-4	25601	Zhipu AI	Proprietary



API Fee

	Gemini (Gemini-1.5-pro)	OpenAI (gpt-4o)	Claude (Claude 3.5 Sonnet)	Hugging Face
Free quota	<ul style="list-style-type: none"> 2 RPM (requests per minute) 32,000 TPM (tokens per minute) 50 RPD (requests per day) 	No	No	FREE
Price Page Link	Link	Link	Link	-
Input token Price (Every 1M tokens)	1.25 2.50 (longer than 128k tokens)	2.50	3	-
Prompt Caching	0.3125 0.625 (longer than 128k tokens) 4.5 per hour for life	1.25*	3.75 (write) 0.30 (read)*	-
Output token Price (Every 1M tokens)	5.00 10.00 (longer than 128k tokens)	10.00	15	-

*life: 5 minutes of inactivity (automatic)



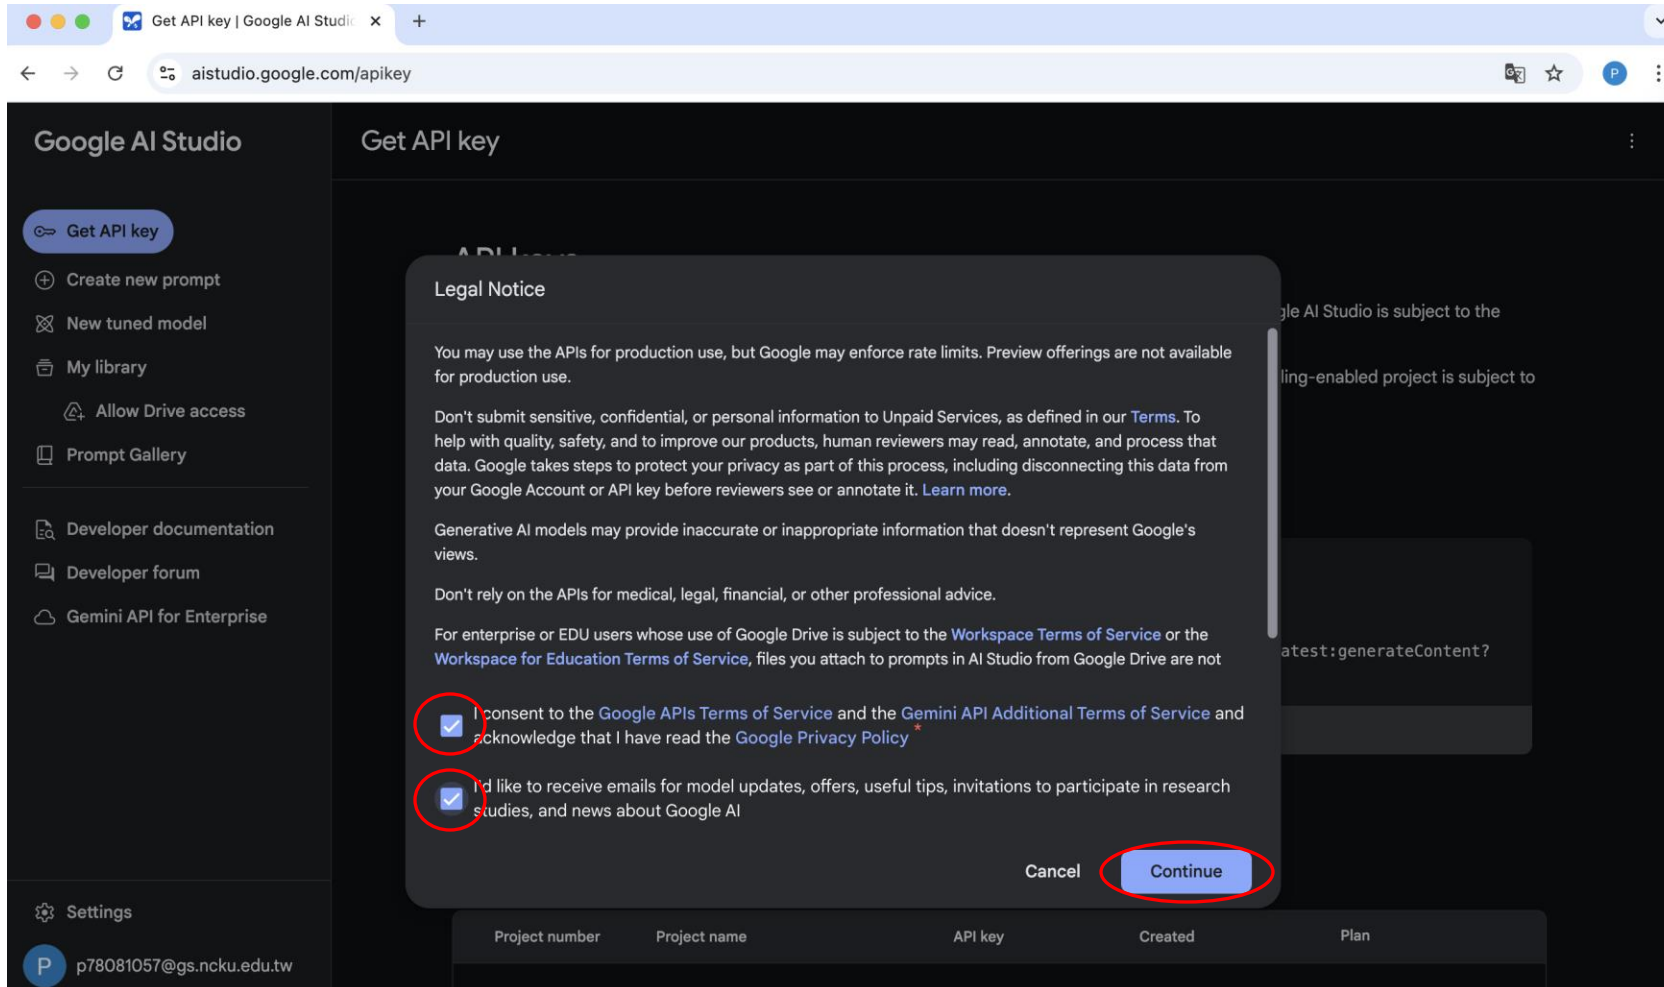
Get Gemini API Key (1)

Go to <https://aistudio.google.com/apikey>

And Log in your Google account

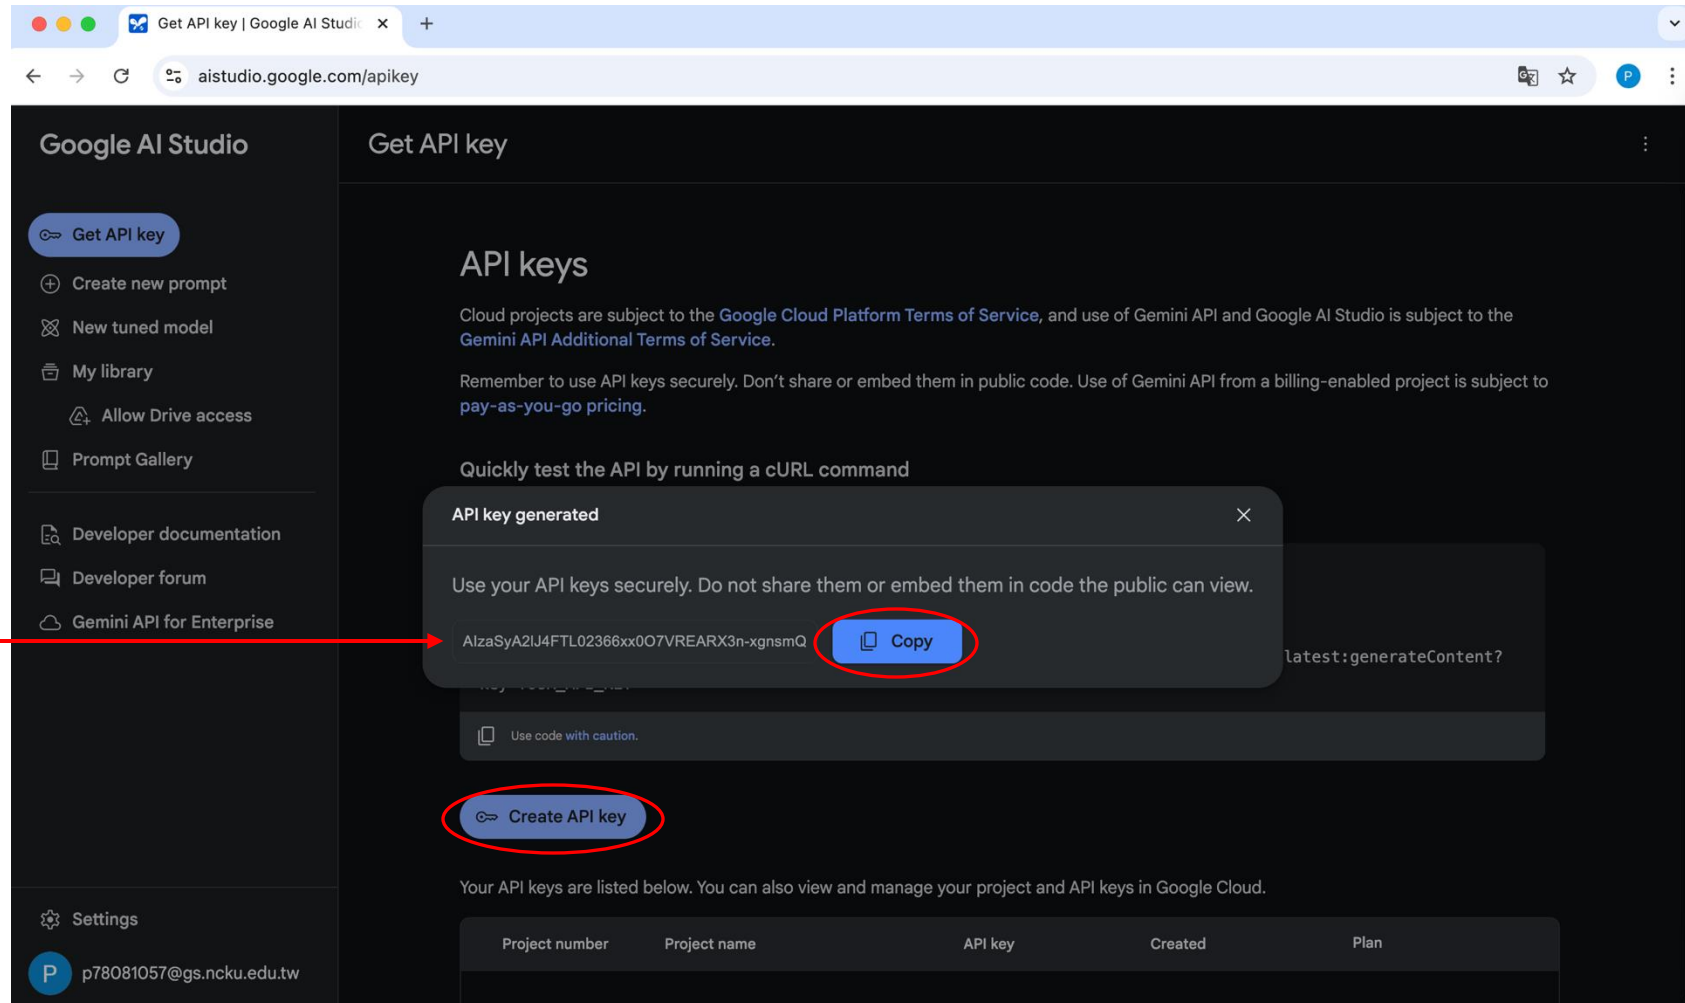


Get Gemini API Key (2)



Get Gemini API Key (3)

Your
API Key



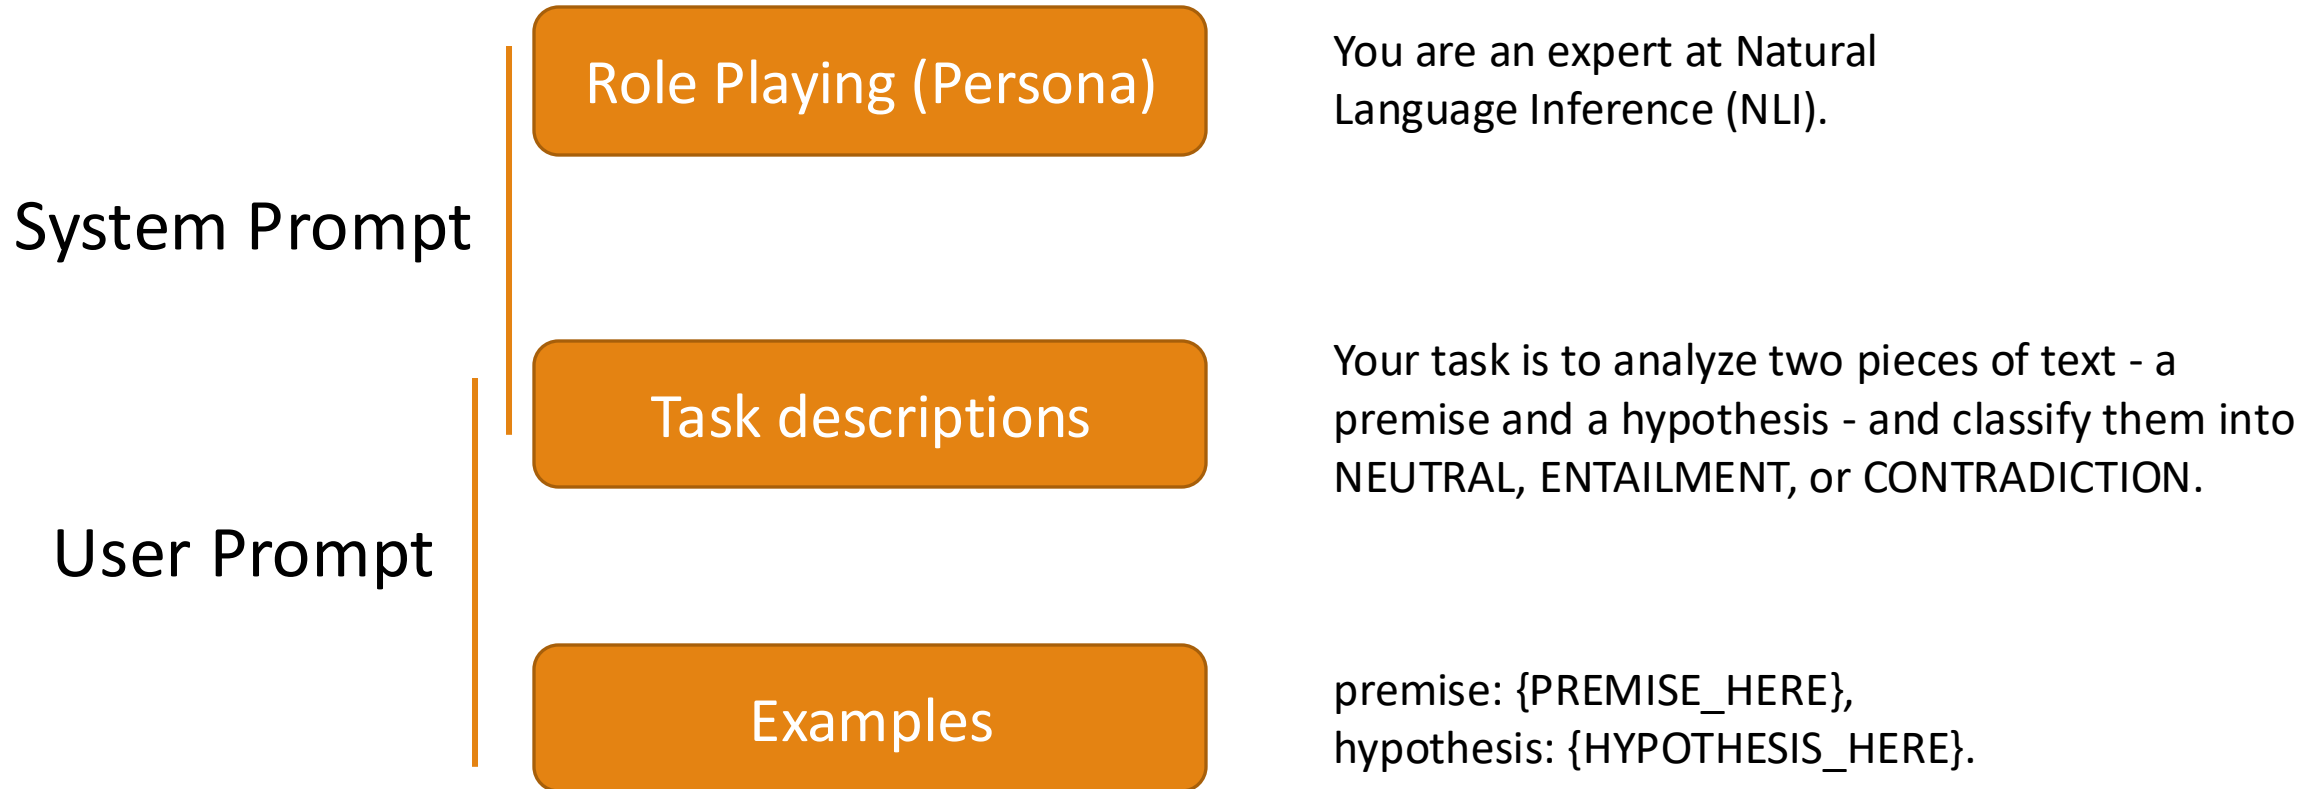
Installation

```
!pip install google-ai-generativelanguage==0.8.3  
!pip install anthropic==0.39.0  
!pip install openai==1.54.5
```

The file for Prompts (prompts.yaml)

```
prompts.yaml
1  system:
2    general: "You are an expert at Natural Language Inference (NLI). Your task is to analyze two pieces of text
3  user:
4    general: "premise: {PREMISE_HERE}, hypothesis: {HYPOTHESIS_HERE}."
5    json_mode: "Generate the classification result in JSON with the key: 'result': str (NEUTRAL, ENTAILMENT, or
6    few: "USER: premise: {PREMISE_1}, hypothesis: {HYPOTHESIS_1}.\nMODEL: {RESULT_1}\nUSER: premise: {PREMISE_2
7  mutual:
8    few_hint: "Please first check the following examples."
```

System Prompt and User Prompt

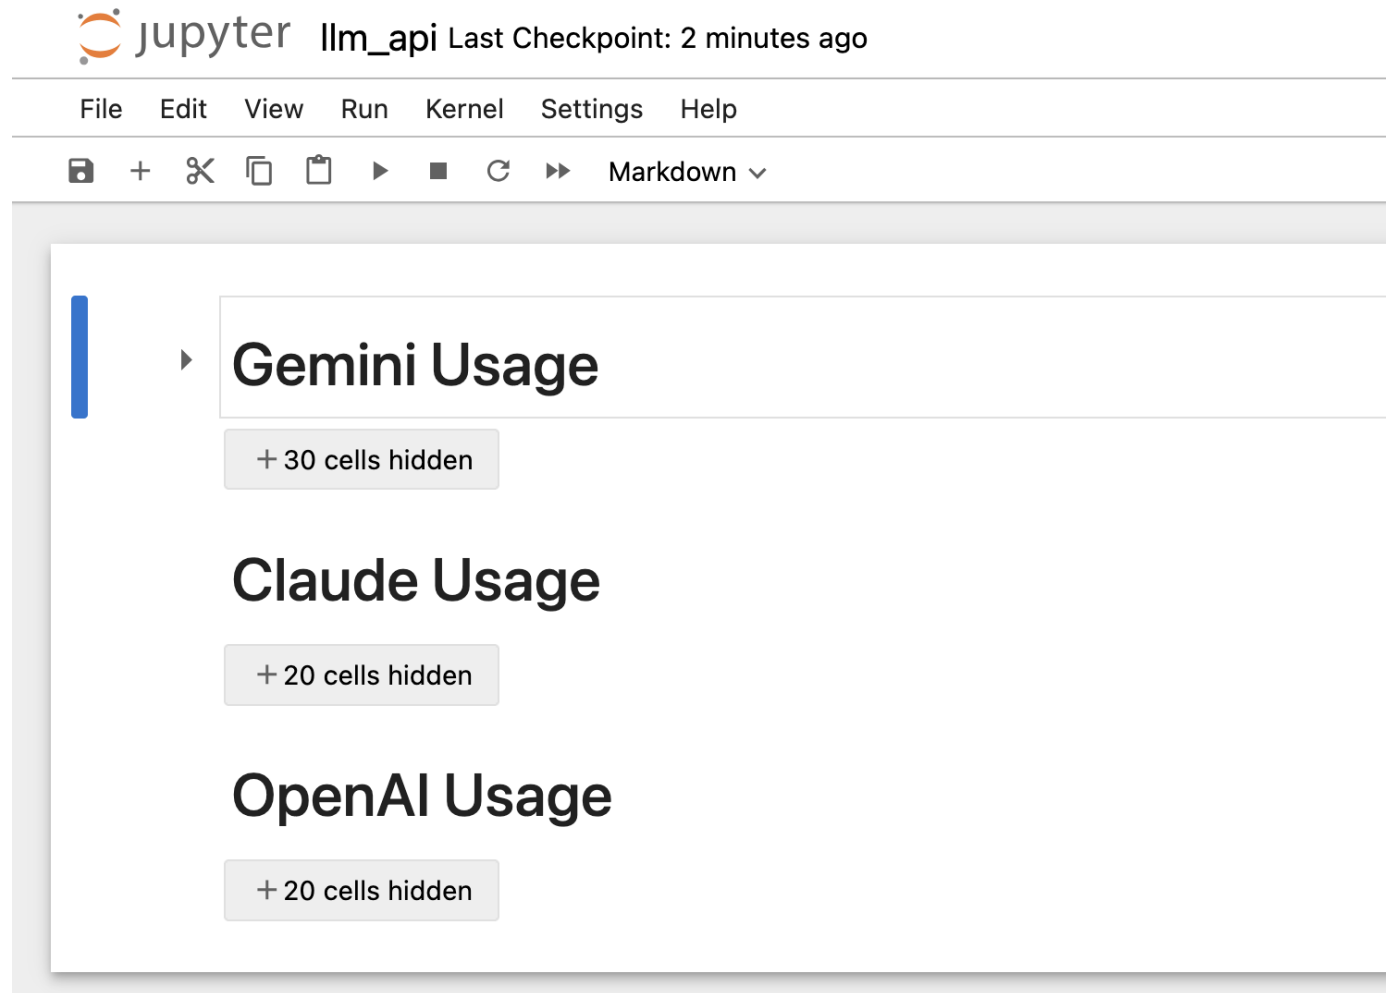


<https://platform.openai.com/docs/guides/prompt-engineering#tactics>



Notebook Executions

Each of the three
parts can be
executed individually.



The screenshot shows a Jupyter Notebook interface with the following elements:

- Header:** jupyter llm_api Last Checkpoint: 2 minutes ago
- Menu Bar:** File Edit View Run Kernel Settings Help
- Toolbar:** Save, New, Copy, Paste, Run, Stop, Refresh, and a dropdown menu for Markdown.
- Content Area:** A list of notebook sections:
 - Gemini Usage**: + 30 cells hidden
 - Claude Usage**: + 20 cells hidden
 - OpenAI Usage**: + 20 cells hidden

Example for Classification

- Dataset: SemEval 2014 Task1-2 (3-class classification)

```
inputs = {  
    "premise": "A group of kids is playing in a yard and an old man  
is standing in the background",  
    "hypothesis": "A group of boys in a yard is playing and a man is  
standing in the background",  
}
```

Setup (Gemini)

<https://ai.google.dev/gemini-api/docs/models/gemini?hl=zh-tw#gemini-1.5-pro>

```
1 import google.generativeai as genai
2 import json
3 from utils import load_prompts
4 prompts = load_prompts("prompts.yaml")
```

```
1 MY_GOOGLE_API_KEY = "MY_GOOGLE_API_KEY"
2 genai.configure(api_key=MY_GOOGLE_API_KEY)
```

```
1 MODEL_NAME = "gemini-1.5-pro"
2 TEMPERATURE = 0
```



The file for Prompts (prompts.yaml)

We use the yaml package to load the “prompts.yaml” into a **Python dictionary**.

```
print(prompts)
```

```
{'system': {'general': 'You are an expert at Natural Language Inference (NLI). Your task is to analyze two pieces of text – a premise and a hypothesis – and classify them into NEUTRAL, ENTAILMENT, or CONTRADICTION.'}, 'user': {'general': 'premise: {PREMISE_HERE}, hypothesis: {HYPOTHESIS_HERE}.', 'json_mode': "Generate the classification result in JSON with the key: 'result': str (NEUTRAL, ENTAILMENT, or CONTRADICTION)", 'few': 'USER: premise: {PREMISE_1}, hypothesis: {HYPOTHESIS_1}.\nMODEL: {RESULT_1}\nUSER: premise: {PREMISE_2}, hypothesis: {HYPOTHESIS_2}.\nMODEL: {RESULT_2}\nUSER: premise: {PREMISE_3}, hypothesis: {HYPOTHESIS_3}.\nMODEL: '}, 'mutual': {'few_hint': 'Please first check the following examples.'}}
```



Basic Usage (Gemini)

```
1 model = genai.GenerativeModel(  
2     MODEL_NAME,  
3     generation_config={"temperature": TEMPERATURE},  
4     system_instruction=system_prompt,  
5 )
```

Set up the Gemini model

```
1 cur_user_prompt = user_prompt.format(  
2     PREMISE_HERE=inputs["premise"],  
3     HYPOTHESIS_HERE=inputs["hypothesis"]  
4 )  
5 print(cur_user_prompt)
```

Format the inputs

premise: A group of kids is playing in a yard and an old man is standing in the background, hypothesis: A group of boys in a yard is playing and a man is standing in the background.

```
1 response = model.generate_content(cur_user_prompt)
```

Run generation



Output of Classification (Gemini)

```
print(response.text)
```

NEUTRAL. The premise states "kids," which could include girls. The hypothesis specifies "boys." The premise says "old man," while the hypothesis just says "man." While an old man is a man, the hypothesis doesn't preclude the man being young or middle-aged. Therefore, the hypothesis is more specific in one way and more general in another, making the relationship neutral.

Count tokens (Gemini)

```
1 num_sys_tokens = str(model.count_tokens(system_prompt)).split(": ")[1].strip()
2 num_usr_tokens = str(model.count_tokens(system_prompt)).split(": ")[1].strip()
3 num_input_tokens = int(num_sys_tokens) + int(num_usr_tokens)
4 print(f"Num of input tokens: {num_input_tokens}")
```

Num of input tokens: 176

```
1 num_output_tokens = str(model.count_tokens(response.text)).split(": ")[1].strip()
2 print(f"Num of output tokens {num_output_tokens}")
```

Num of output tokens 128



Output of Classification (Gemini)

```
print(response.text)
```

NEUTRAL. The premise states "kids," which could include girls. The hypothesis specifies "boys." The premise says "old man," while the hypothesis just says "man." While an old man is a man, the hypothesis doesn't preclude the man being young or middle-aged. Therefore, the hypothesis is more specific in one way and more general in another, making the relationship neutral.

How to get structured output for evaluating model performance?

JSON mode -> {"result": "NEUTRAL"}

Generate structured output in JSON (Gemini)

```
1 user_prompt_json = prompts["user"]["general"] + " " + prompts["user"]["json_mode"]
```

```
1 cur_user_prompt = user_prompt_json.format(  
2     PREMISE_HERE=inputs["premise"],  
3     HYPOTHESIS_HERE=inputs["hypothesis"]  
4 )  
5 print(cur_user_prompt)
```

premise: A group of kids is playing in a yard and an old man is standing in the background, hypothesis: A group of boys in a yard is playing and a man is standing in the background. Generate the classification result in JSON with the key: 'result': str (NEUTRAL, ENTAILMENT, or CONTRADICTION)

```
1 response = model.generate_content(  
2     cur_user_prompt,  
3     generation_config={  
4         "response_mime_type": "application/json",  
5     },  
6 )
```

Set up the config for Gemini to output in JSON

<https://ai.google.dev/gemini-api/docs/structured-output?lang=python>



Output of Classification in JSON (Gemini)

```
1 result_json = json.loads(response.text)
2 print(type(result_json))
3 print(result_json)
```

```
<class 'dict'>
{'result': 'NEUTRAL'}
```

Few-shot Prompts (Gemini, [ref](#))

- Few-shot: some examples with labels are provided to an LLM
- This approach already prompts the model to follow the output format (JSON mode may not be needed.)

```
1 fs_user_prompt = prompts["user"]["few"]
```

```
1 cur_fs_user_prompt = fs_user_prompt.format(  
2     PREMISE_1="A group of kids is playing in a yard and an old man is standing in the background",  
3     HYPOTHESIS_1="A group of boys in a yard is playing and a man is standing in the background",  
4     RESULT_1="{ 'result': 'NEUTRAL' }",  
5     PREMISE_2="A man, a woman and two girls are walking on the beach",  
6     HYPOTHESIS_2="A group of people is on a beach",  
7     RESULT_2="{ 'result': 'ENTAILMENT' }",  
8     PREMISE_3="Two teams are competing in a football match",  
9     HYPOTHESIS_3="Two groups of people are playing football",  
10 )  
11 print(cur_fs_user_prompt)
```

USER: premise: A group of kids is playing in a yard and an old man is standing in the background, hypothesis: A group of boy
s in a yard is playing and a man is standing in the background.

MODEL: { 'result': 'NEUTRAL' }

USER: premise: A man, a woman and two girls are walking on the beach, hypothesis: A group of people is on a beach.

MODEL: { 'result': 'ENTAILMENT' }

USER: premise: Two teams are competing in a football match, hypothesis: Two groups of people are playing football.

MODEL:



The file for Prompts (prompts.yaml)

We use the yaml package to load the “prompts.yaml” into a **Python dictionary**.

```
print(prompts)
```

```
{'system': {'general': 'You are an expert at Natural Language Inference (NLI). Your task is to analyze two pieces of text – a premise and a hypothesis – and classify them into NEUTRAL, ENTAILMENT, or CONTRADICTION.'}, 'user': {'general': 'premise: {PREMISE_HERE}, hypothesis: {HYPOTHESIS_HERE}.', 'json_mode': "Generate the classification result in JSON with the key: 'result': str (NEUTRAL, ENTAILMENT, or CONTRADICTION)", 'few': 'USER: premise: {PREMISE_1}, hypothesis: {HYPOTHESIS_1}.\nMODEL: {RESULT_1}\nUSER: premise: {PREMISE_2}, hypothesis: {HYPOTHESIS_2}.\nMODEL: {RESULT_2}\nUSER: premise: {PREMISE_3}, hypothesis: {HYPOTHESIS_3}.\nMODEL:'}, 'mutual': {'few_hint': 'Please first check the following examples.'}}
```



Few-shot Prompting (Gemini)

```
1 fs_system_prompt = prompts["system"]["general"] + " " + prompts["mutual"]["few_hint"]
2 print(fs_system_prompt)
```

You are an expert at Natural Language Inference (NLI). Your task is to analyze two pieces of text – a premise and a hypothesis – and classify them into NEUTRAL, ENTAILMENT, or CONTRADICTION. Please first check the following examples.

```
1 model = genai.GenerativeModel(
2     MODEL_NAME,
3     generation_config={"temperature": TEMPERATURE},
4     system_instruction=fs_system_prompt,
5 )
```

```
1 response = model.generate_content(
2     cur_fs_user_prompt,
3     generation_config={
4         "response_mime_type": "application/json",
5     },
6 )
```

```
1 result_json = json.loads(response.text)
2 print(result_json)
```

```
{'result': 'ENTAILMENT'}
```



Example for Summarization (Gemini)

- Dataset: LCSTS (Chinese Abstractive Summarization)

```
1 system_prompt = """你是個中文文本摘要的專家，現在請你對一篇輸入的文章進行摘要。"""
2
3 model = genai.GenerativeModel(
4     MODEL_NAME,
5     generation_config={"temperature": TEMPERATURE},
6     system_instruction=system_prompt,
7 )
```

```
1 input_source_txt = "新华社授权于18日全文播发修改后的《中华人民共和国立法法》，修改后的立法法分为“总则”“法律”“行政法规”“地方性法规、自治条例和单行条例”“法律解释”“立法监督”“附则”七章，共105条。"
```

```
1 response = model.generate_content(input_source_txt)
```

```
1 print(response.text)
```

'新修订的《中华人民共和国立法法》已正式发布，共六章105条，涵盖总则、法律、行政法规、地方性法规及规章、适用与备案审查以及附则等方面。\\n'



Prompt Caching

<https://ai.google.dev/gemini-api/docs/caching?lang=python>

```
# Create a cache with a 5 minute TTL
cache = caching.CachedContent.create(
    model='models/gemini-1.5-flash-001',
    display_name='sherlock jr movie', # used to identify the cache
    system_instruction=(
        'You are an expert video analyzer, and your job is to answer '
        'the user\'s query based on the video file you have access to.'
    ),
    contents=[video_file],
    ttl=datetime.timedelta(minutes=5),
)

# Construct a GenerativeModel which uses the created cache.
model = genai.GenerativeModel.from_cached_content(cached_content=cache)

# Query the model
response = model.generate_content([
    'Introduce different characters in the movie by describing '
    'their personality, looks, and names. Also list the timestamps '
    'they were introduced for the first time.'])
```

When to use?

- Chatbots with extensive system instructions
- Queries against large document sets
- Analysis of a lengthy video

A very big (long) file

System instruction and the video file is cached.
Cached tokens are in a lower price.



Notebook Executions

Most of them
are similar.

jupyter llm_api Last Checkpoint: 2 minutes ago

File Edit View Run Kernel Settings Help

Save + Delete Copy Paste Run Interrupt Refresh Markdown ▾

▶ **Gemini Usage**
+ 30 cells hidden

Claude Usage
+ 20 cells hidden

OpenAI Usage
+ 20 cells hidden

API Keys

- OpenAI API
 - <https://platform.openai.com/api-keys>
- Claude API
 - <https://console.anthropic.com/settings/keys>



Difference in Few-shot Prompting

OpenAI API ([ref](#))

```
1 messages
```

```
[{'role': 'system',
  'content': 'You are an expert at Natural Language Inference (NLI). Your task is to analyze two pieces of text - a premise and a hypothesis - and classify them into NEUTRAL, ENTAILMENT, or CONTRADICTION. Please first check the following examples.'},
 {'role': 'user',
  'content': "premise: A group of kids is playing in a yard and an old man is standing in the background, hypothesis: A group of boys in a yard is playing and a man is standing in the background. Generate the classification result in JSON with the key: 'result': str (NEUTRAL, ENTAILMENT, or CONTRADICTION)"},
 {'role': 'assistant', 'content': '{"result': NEUTRAL}"},
 {'role': 'user',
  'content': "premise: A man, a woman and two girls are walking on the beach, hypothesis: A group of people is on a beach. Generate the classification result in JSON with the key: 'result': str (NEUTRAL, ENTAILMENT, or CONTRADICTION)"},
 {'role': 'assistant', 'content': '{"result': ENTAILMENT}"},
 {'role': 'user',
  'content': "premise: Two teams are competing in a football match, hypothesis: Two groups of people are playing football. Generate the classification result in JSON with the key: 'result': str (NEUTRAL, ENTAILMENT, or CONTRADICTION)"}]
```

```
1 response = client.chat.completions.create(
2     model=MODEL_NAME,
3     response_format={"type": "json_object"},
4     messages=messages, ← list
5     temperature=TEMPERATURE,
6     # max_tokens=max_tokens,
7 )
8 result_json = json.loads(response.choices[0].message.content)
```



Difference in Few-shot Prompting

Claude API
([ref](#))

```
1 cur_fs_user_prompt = fs_user_prompt.format(  
2     PREMISE_1="A group of kids is playing in a yard and an old man is standing in the background",  
3     HYPOTHESIS_1="A group of boys in a yard is playing and a man is standing in the background",  
4     RESULT_1="{ 'result': 'NEUTRAL' }",  
5     PREMISE_2="A man, a woman and two girls are walking on the beach",  
6     HYPOTHESIS_2="A group of people is on a beach",  
7     RESULT_2="{ 'result': 'ENTAILMENT' }",  
8     PREMISE_3="Two teams are competing in a football match",  
9     HYPOTHESIS_3="Two groups of people are playing football",  
10 )  
11 print(cur_fs_user_prompt)
```

Same as Gemini

USER: premise: A group of kids is playing in a yard and an old man is standing in the background, hypothesis: A group of boys in a yard is laying and a man is standing in the background.
MODEL: {'result': 'NEUTRAL'}
USER: premise: A man, a woman and two girls are walking on the beach, hypothesis: A group of people is on a beach.
MODEL: {'result': 'ENTAILMENT'}
USER: premise: Two teams are competing in a football match, hypothesis: Two groups of people are playing football.
MODEL:

```
1 response = client.messages.create(  
2     model=MODEL_NAME,  
3     max_tokens=1000,  
4     temperature=TEMPERATURE,  
5     system=fs_system_prompt,  
6     messages=[  
7         {  
8             "role": "user",  
9             "content": [{"type": "text", "text": cur_user_prompt}],  
10        },  
11    ],  
12 )  
13 result = response.content[0].text
```

↑
string



Count the number of tokens

OpenAI API

```
print(f"Num of input tokens {response.usage.prompt_tokens}")  
print(f"Num of output tokens {response.usage.completion_tokens}")
```

Claude API

```
print(f"Num of input tokens {response.usage.input_tokens}")  
print(f"Num of output tokens {response.usage.output_tokens}")
```

Count the number of tokens (OpenAI)

OpenAI Tokenizer:

<https://platform.openai.com/tokenizer>

Question: How to calculate the number of tokens for Claude and Gemini?

- All the three APIs (OpenAI, Anthropic, and Gemini) can calculate number of tokens after processing a request.
- Only **OpenAI** offers **pre-calculator** to count tokens.

Learn about language model tokenization

OpenAI's large language models process text using **tokens**, which are common sequences of characters found in a set of text. The models learn to understand the statistical relationships between these tokens, and excel at producing the next token in a sequence of tokens. [Learn more.](#)

You can use the tool below to understand how a piece of text might be tokenized by a language model, and the total count of tokens in that piece of text.

GPT-4o & GPT-4o mini

GPT-3.5 & GPT-4

GPT-3 (Legacy)

Explain to me how AI works

Clear

Show example

Tokens

6

Characters

26

Explain to me how AI works



Further Learning

- Prompt caching:
 - [Practical Implementations using Gemini API](#)
 - [OpenAI API](#)
 - [Claude API \(beta\)](#)
- Batch API (Batch predictions):
 - [OpenAI API](#)
 - [Claude API \(beta\)](#)



Thank you!

